

Zoha Rakotomalala

AI/ML SOFTWARE ENGINEER · CLOUD & MLOPS

Paris, France

✉ zoha.rakotomalala@gmail.com | 🌐 zoha-rakotomalala | 🌐 ZohaRak

ESA Recruitment Team, Data & AI Unit

July 1, 2026

EUROPEAN SPACE AGENCY

DIGITAL, DATA AND INFORMATION TECHNOLOGY DEPARTMENT

PARIS, FRANCE

Application: Junior Professional in AI/ML Software Engineering (Req. 20679)

Dear ESA Recruitment Team,

I am applying for the Junior Professional position in AI/ML Software Engineering because it sits exactly where I want to grow. For two years I have built and operated reliable software and cloud infrastructure in production; alongside that, I have become the engineer on my team who adopts, evaluates and champions AI tooling. What I have not yet done, and what this role is built to develop, is take machine-learning models themselves from prototype to dependable operational use. The software-engineering rigour, the cloud and CI/CD foundation, and the genuine pull toward AI are already there; the Junior Professional Programme is the structured runway to turn that into ML/MLOps depth.

At Amazon I own operational excellence and infrastructure modernisation for a team running cross-marketplace pricing services. I recently delivered a cross-region cloud migration and a Redshift descale that removed roughly 15% of a cluster's footprint for about €135K/year in recurring savings; I build and deploy infrastructure-as-code (AWS CDK) on Linux, and I hold oncall responsibility for production systems where reliability is not negotiable. Before that I spent nearly two years as a full-stack engineer (Java, TypeScript) on price-transparency and legal-compliance features for European marketplaces, where I owned a market launch end to end via A/B experiment, and authored the post-mortem after a compliance incident, driving the mitigations to closure. I mention the incident deliberately: the part of your mandate about ensuring AI solutions comply with governance and security constraints is not abstract to me. I have lived the cost of getting compliance wrong in production, and the rigour that came from fixing it.

My path to engineering was not the standard one, and I think it is the reason I fit this role rather than a caveat to it. I came through business school (EDHEC, then a year at UC Berkeley) on the maths and computer-science foundation I laid in prépa at Lycée Saint-Louis, took every CS and data course I could reach, and taught myself the rest. My first engineering home was Dataiku, an ML-platform company, where I shipped a 1,900-line feature indexing database relationships into Elasticsearch and rendering them as an interactive graph. Being largely self-taught has made me quick to close gaps, comfortable with ambiguity, and drawn to problems that do not yet have an obvious answer. That is the mindset I would bring to learning the ML-serving stack.

The daily tools of this role are ones I already use: Python, Linux, Git and CI/CD. The mission is where I want to build depth: the model lifecycle, covering deployment, monitoring, a model catalogue, and an LLM inference service across on-prem and cloud. My AI experience so far is as a hands-on user and advocate rather than a builder of models: I pioneered my team's adoption of AI-assisted development tooling and presented on the Model Context Protocol and its architecture for AI-integrated systems to over forty engineers at an internal conference. I am ready to move from using these systems to engineering and operating them, and the Programme's four-year arc with its mobility component is exactly the structured runway to do it.

I am a French national based in Paris, fluent in French and English, and available to start within your timelines. I would welcome the opportunity to bring my engineering foundation to ESA and grow into helping turn its AI work into dependable operational capability.

With respect and enthusiasm,

Zoha Rakotomalala